Weidong Tang 汤维栋

Incoming PhD · CRO (Founding Member), Cardinal AI Lab, InfRec

wdtang0705@gmail.com · wd7ang.github.io · GitHub · Google Scholar · WeChat: d7ang0705

EDUCATION

Xidian University — Sep 2023 - Jul 2027 (expected)
B.E., School of Telecommunication Engineering · GPA: 3.8/4.0
Advised by Prof. Xiumei Wang.

EXPERIENCE

Cardinal AI Lab, InfRec — Jul 2026 - Present
Chief Research Officer (Founding Member)
- Lead research on recursive self-improvement and long-horizon world interaction, enabling AI agents to retain knowledge and control embodied robots.
- Guide projects from problem formulation to prototype development, working with engineering teams on practical deployment.

Alibaba — May 2026 - Jul 2026
MLLM Research Intern
- Researched post-training for multimodal foundation models, incorporating e-commerce domain knowledge for product retrieval and similar-item search.
- Developed training and evaluation pipelines for product representation learning, visual-semantic alignment, and retrieval in large-scale e-commerce settings.

TeleAI — Jan 2026 - Mar 2026
LLM Research Intern
- Worked on AI infrastructure for LLM training and inference on Huawei Ascend 910B, including precision alignment.

National University of Singapore · HPC-AI Lab — Oct 2024 - Oct 2025
Research Assistant
Advised by Prof. Yang You, Prof. Wangbo Zhao, and Dr. Pengfei Zhou.
- Conducted research on efficient machine learning and the evaluation of multimodal models at the HPC-AI Lab.
- Led projects from initial ideas through publication, resulting in three papers accepted at CVPR and ACL.

PUBLICATIONS

* Equal contribution · † Corresponding author · ‡ Project leader

01 **The Trinity: Capability-Driven Recursive Self-Improvement for Multimodal Agents**
W. Tang*, R. Wang*, J. Zhang*, H. Wang, H. Chen, Z. Zhang, Y. Ma, Z. Jiang, Z. Guo, H. Li, Z. Jiao, Z. Cui, X. Yin, S. Liang, Z. Wan, H. Lin, W. Zhao, K. Zhang, P. Zhou†, O. Zhang†, & Y. You†.
arXiv preprint, 2026 Paper · GitHub · Project Page

02 **GroupToM-Bench: Benchmarking Group Theory of Mind and Nonlinear Social Emergence in MLLMs**
W. Tang, J. Li, Y. Hou, Z. Mei, C. Zhang, X. Wan, Z. Liang, P. Zhou†, Y. You, & W. Zhao†.
ACL 2026 (19% acceptance rate, CCF-A top-tier conference) Paper · Dataset

03 **MagicBench: Diagnosing Visual Agency Loss and Semantic Dependency in Multimodal LLMs**
T. D. Huang, W. Tang‡, W. Q. Xu, & X. Guo†.
ACL 2026 (19% acceptance rate, CCF-A top-tier conference) Paper · GitHub

04 **Efficient Video Object Segmentation and Tracking with Recurrent Dynamic Submodel**
W. Tang, Z. Liang, X. Wan, C. Zhu, Z. Xu, P. Zhou†, Y. Song, Y. You, & W. Zhao†.
CVPR 2026 (25% acceptance rate, CCF-A top-tier conference) Paper

05 **Inference-Time Scaling for Visual AutoRegressive Modeling by Searching Representative Samples**
W. Tang, X. Wan, S. Li, & X. Wang†.
PRCV 2025 (29% acceptance rate, CCF-C) Paper · GitHub

06 **EarthVerse: Benchmarking Scientific Agents Across Dynamic Earth Systems and Natural Hazards**
Z. Cui, X. Yin, Y. Tang, X. Zhang, Y. Hu, S. Zhong, **W. Tang**, Y. Liang, W. Li, M. Jin, S. Pan, Y. Kang, D. Zhuang, & J. Zhao.
arXiv preprint, 2026 Paper · GitHub · Dataset · Project Page

07 **IndustryBench-MIPU: Benchmarking Multi-Image Attribute Value Extraction for Industrial Products**
H. Qi, J. Cao, Y. Zhang, X. Wang, **W. Tang**, B. Chen, C. Huo, H. Pan, H. You, J. Li, Y. Wang, & L. Ding.
arXiv preprint, 2026 Paper · GitHub · Dataset

08 **MDK12-Bench: A Multi-Discipline Benchmark for Evaluating Reasoning in Multimodal Large Language Models**
P. Zhou, X. Peng, F. Zhang, Z. Xu, J. Ai, Y. Qiu, W. Zhao, J. Song, C. Li, **W. Tang**, Z. Li, H. Zhang, Z. Li, X. Mao, Y. Feng, J. Sun, K. Wang, X. Chang, W. Shao, Y. You, & K. Zhang†.
AAAI 2026 (17% acceptance rate, CCF-A) Paper · GitHub

09 **EA-ViT: Efficient Adaptation for Elastic Vision Transformer**
C. Zhu, W. Zhao†, H. Zhang, Y. Zhou, **W. Tang**, S. Wang, Z. Yuan, Y. Shang, X. Peng, K. Wang, & D. Yang†.
ICCV 2025 (25% acceptance rate, CCF-A top-tier conference) Paper · GitHub

10 **Pebench: A Fictitious Dataset to Benchmark Machine Unlearning for Multimodal Large Language Models**
Z. Xu, P. Zhou, **W. Tang**, J. Ai, W. Zhao, K. Wang, X. Peng, W. Shao, H. Yao, & K. Zhang†.
arXiv preprint, 2025 Paper · Dataset · Project Page

11 **IEBench: Benchmarking Unlearning Identity-Event Entangled Concepts for MLLM-Based Systems in Agentic Edge Networks**
Z. Xu, P. Zhou, **W. Tang**, J. Ai, W. Zhao, K. Wang, W. Shao, X. Peng, S. Zhao, H. Yao, & K. Zhang†.
IEEE Transactions on Network Science and Engineering, 2026 (IF=7.9, JCR-Q1) Paper

SELECTED PROJECTS

WowPage
Weidong Tang, Yue Su.
A clean and standalone homepage template. GitHub · RedNote
Developed an open-source, lightweight personal homepage framework featuring customizable layouts for researchers to present publications, projects, and academic profiles.

AWARDS

Academic Scholarship, Xidian University (Top 5%) 2026
Academic Scholarship, Xidian University (Top 5%) 2025

PROFESSIONAL SERVICES

AI Team Lead, Xidian-Inspur Club Jun 2025 - Present
Coordinator, Xidian Graduate Admissions Experience-Sharing Group (3,000 members) 2026